

Python Code Documentation

Function

`admissible_subsets`

Description

Performs variable selection in the linear model setting. Can to handle both $p \leq n$ and $p > n$ situations. See the reference for further details.

Usage

```
admissible_subsets(y, X, N, steps, burnin, po, PropWeights)
```

Arguments

- y:** An $n \times 1$ NumPy array of response data.
- X:** An $n \times p$ NumPy array of predictors.
- N:** The number of importance samples used to estimate $E(h(\beta_M))$ within the pseudo-marginal MCMC. Default is 100.
- steps:** The number of MCMC steps.
- burnin:** The number of initial *steps* to discard.
- po:** Belief about the number of predictors in the true model. Can be chosen via cross-validation with `admissible_subsets_cv`.
- PropWeights:** The weights used for proposing which predictors to add or drop as the MCMC samples index sets $M \subset \{1, \dots, p\}$. Default is to use squared coefficient estimates from elastic net, via the `ElasticNetCV` function from the ‘sklearn.linear_model’ Python module, added by n^{-2} .

Values

- chain:** A $(steps - burnin) \times p$ NumPy array containing the MCMC sample path (or trace) over index sets M , after *burnin* number of steps.
- postSample:** A NumPy array containing the indices (in each row) for every M visited in the MCMC sample path.
- postProbs:** A list containing the relative frequencies for which each of the index sets M in *postSample* was visited in the MCMC sample path.

AcceptRatio: The number of MCMC steps in which a proposed index set M was accepted.

References

J. P. Williams and J. Hannig (2017). Non-penalized variable selection in high-dimensional linear model settings via generalized fiducial inference. *Submitted*.

Supplement: https://jonathanpw.github.io/assets/WilliamsHannig_supplement.pdf

Function

`admissible_subsets_cv`

Description

Performs k -fold cross-validation to choose po in `admissible_subsets`, using an implementation of `admissible_subsets`. See the reference for further details.

Usage

`admissible_subsets_cv(y, X, N, steps, burnin, grid, num_folds)`

Arguments

y:	An $n \times 1$ NumPy array of response data.
X:	An $n \times p$ NumPy array of predictors.
N:	The number of importance samples used to estimate $E(h(\beta_M))$ within the pseudo-marginal MCMC, during the cross-validation procedure. Default is 30.
steps:	The number of MCMC steps, during the cross-validation procedure. Default is 200.
burnin:	The number of initial <i>steps</i> to discard, during the cross-validation procedure. Default is 100.
grid:	A list of po values to perform cross-validation over. Default is $[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]$.

num.folds: The number of folds to partition the data into for performing k -fold cross-validation. Default is 10.

Values

po: The optimal tuning parameter based on the implemented cross-validation.

References

J. P. Williams and J. Hannig (2017). Non-penalized variable selection in high-dimensional linear model settings via generalized fiducial inference. *Submitted*.

Supplement: https://jonathanpw.github.io/assets/WilliamsHannig_supplement.pdf